



Bahir Dar University
Faculty of Computing
Department of Software Engineering

OSSP Individual Assignment

System Call Implementation - `execve()`

Full Name: Elias Zeleke

Student ID: BDU1601353

Section: B

System Call: `execve()`

Submitted To: Lecturer Wendimu Baye

Submission Date: 16/08/2017 E.C

System Call Implementation

❖ `execve()` in User Space

The `execve ()` system call in C allows a program to execute a new process by replacing the current one. In user space, this system call is accessed via a glibc wrapper, which internally uses the `syscall ()` function to request the kernel's implementation of `execve`.

This wrapper handles communication with the kernel and passes the necessary arguments: the path to the executable, the arguments for the new program, and the environment variables.

User-Space Code:

```
#include <unistd.h>

#include <sys/syscall.h>

#include <errno.h>

int execve(const char *pathname, char *const argv[], char *const envp[]) {

return syscall(SYS_execve, pathname, argv, envp);

}
```

this code defines how the high-level `execve()` call in C is translated into a low-level system call using `syscall()`. The constant `SYS_execve` is defined in system headers and maps to the correct syscall number for your platform.

❖ Kernel Space: Actual System Call Logic

In the Linux kernel, the `execve()` system call is defined inside the `fs/exec.c` file. The `syscall` is implemented and registered using the `SYSCALL_DEFINE3` macro, which specifies that `execve()` takes three arguments.

Kernel-Level Code:

```
SYSCALL_DEFINE3(execve,
    const char __user *, filename,
    const char __user *const __user *, argv,
    const char __user *const __user *, envp)
{
    return do_execve(getname(filename), argv, envp);
}
```

Key Steps in Kernel Logic:

- `getname(filename)`: Validates and retrieves the binary's filename from user space.
- `do_execve ()`: Responsible for loading the new binary, setting up memory, clearing the current process image, and initializing the new one.
- If successful, the calling process image is destroyed and replaced by the new program.
- If an error occurs, it returns an error code.

This syscall does not return on success — the process context is changed, and the new program takes control.

❖ How `execve()` Works

The `execve()` system call is unique because it replaces the currently running program. Unlike `fork()`, which creates a new child process, `execve()` completely overwrites the current process image with a new executable.

Key Behavior:

- It does not return if the new program loads successfully.
- If execution fails (e.g., due to a missing binary or permission error), `execve()` returns -1 and sets the `errno` variable.
- All previous memory content, file descriptors (unless marked `O_CLOEXEC`), and signal handlers are removed.

This makes `execve()` critical for creating shells, script runners, or OS-level tools that manage programs.

❖ `execve()` Example Code and Output

`execdemo.c` - A Program That Executes `ls -l /home`

This example demonstrates the use of `execve()` to run the `ls` command on the `/home` directory. It replaces the current process with a new one pointing to `/bin/ls`.

```

#include <unistd.h>

#include <stdio.h>

int main() {

    char *program = "/bin/ls";

    char *args[] = { "ls", "-l", "/home", NULL };

    execve(program, args, NULL);

    // If we reach here, execve has failed

    perror("execve failed");

    return 1;

}

```

Compile and Run:

```

gcc execdemo.c -o execdemo
./execdemo

```

Expected Output:

If /home contains files or folders, the output will look like:

```

/home:
total 8
drwxr-xr-x 2 root root 4096 Apr 24 09:59 user1
-rw-r--r-- 1 root root 10 Apr 24 10:00 notes.txt

```

If the binary or path is invalid, the error handler prints:

```

execve failed: No such file or directory

```

❖ [hello.c](#) - A Basic Output Program for Testing GCC

This is a simple program used to verify that your compiler and execution environment (e.g., Tiny Core Linux) are working correctly. It prints a static message to the terminal.

```
#include <stdio.h>

int main() {

printf("This is Operating assignment at Bahir Dar University.\n");

return 0;
}
```

Compile and Run:

```
gcc hello.c -o hello
```

```
./hello
```

Output:

This is Operating assignment at Bahir Dar University.

❖ Under the Hood: Execution Layers

Layer	Component	Role
User Space	<code>execve()</code> (glibc)	Wrapper function that calls <code>syscall()</code>
Syscall Entry	<code>execve()</code> (glibc)	Registers syscall and passes args to kernel logic
Kernel Logic	<code>do_execve()</code>	Loads new executable, maps memory, replaces process image
File System	<code>search_binary_handler()</code>	Locates and validates the binary (e.g., ELF or script)

This layered architecture ensures that the request to run a new program safely flows from user space into the kernel, all the way to the file system.

❖ Summary

`execve()` is a low-level Linux system call used to load and run new programs.

Unlike `fork()`, which creates a child process, `execve()` replaces the current process entirely.

It is essential for building shells, scripting environments, and boot-time process initialization.

This assignment demonstrated both the code-level implementation of `execve()` and its underlying kernel logic using simple, real-world examples.

Tiny Core Linux was used as the platform due to its lightweight, modular nature — making it ideal for system-level programming exploration.

...Thank you...